

forchapter5

Haziq Jamil

18/03/2020

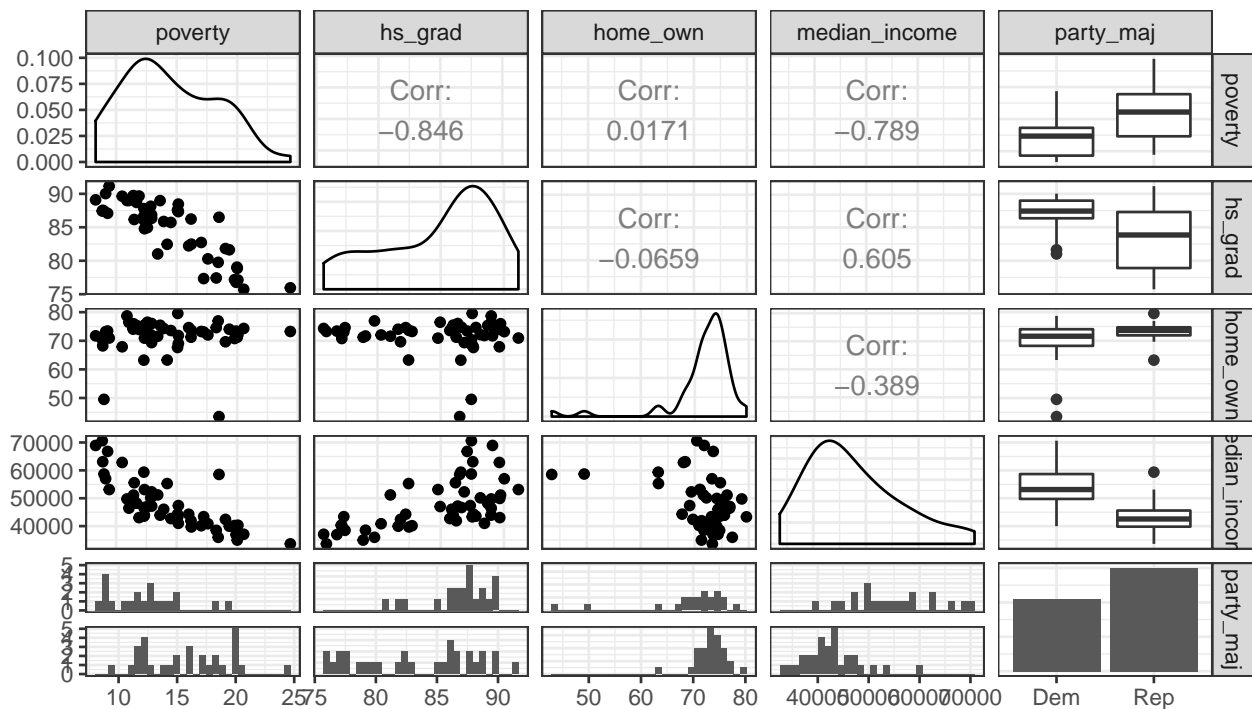
```
## Parsed with column specification:
## cols(
##   state = col_character(),
##   poverty = col_double(),
##   hs_grad = col_double(),
##   home_own = col_double(),
##   median_income = col_double(),
##   party_maj = col_character()
## )
```

poverty

```
## # A tibble: 51 x 6
##   state poverty hs_grad home_own median_income party_maj
##   <chr>      <dbl>   <dbl>   <dbl>      <dbl> <chr>
## 1 Alabama    19.9    76.8    73.4    36963. Rep
## 2 Alaska     12.3    86.6    63.2    59351. Rep
## 3 Arizona     19.0    81.8    69.7    42418. Rep
## 4 Arkansas    20.0    78.9    71.2    34983. Rep
## 5 California  14.2    82.5    63.3    55266. Dem
## 6 Colorado    12.9    88.1    72.1    50136. Dem
## 7 Connecticut  8.31    89.1    71.7    68935. Dem
## 8 Delaware    11.5    86.2    74.7    55568. Dem
## 9 District of Columbia 18.5    86.5    43.5    58526. Dem
## 10 Florida    16.0    82.2    74.6    44269. Rep
## # ... with 41 more rows
```

```
GGally::ggpairs(poverty[, -1]) + theme_bw()
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
mod <- lm(formula = poverty ~ hs_grad + home_own + I(median_income/1000) +
           party_maj, data = poverty)
summary(mod)
```

```
##
## Call:
## lm(formula = poverty ~ hs_grad + home_own + I(median_income/1000) +
##     party_maj, data = poverty)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.8794 -0.8104 -0.0762  0.8412  3.2827
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    76.07336    4.58210   16.602 < 2e-16 ***
## hs_grad        -0.45345    0.05599   -8.098 2.12e-10 ***
## home_own       -0.14614    0.03538   -4.130 0.000151 ***
## I(median_income/1000) -0.26101    0.03427   -7.616 1.10e-09 ***
## party_majRep   -0.51100    0.51347   -0.995 0.324857
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.354 on 46 degrees of freedom
## Multiple R-squared:  0.8859, Adjusted R-squared:  0.876
## F-statistic: 89.28 on 4 and 46 DF,  p-value: < 2.2e-16
```

```
X <- model.matrix(mod); head(X) # n by (p + 1) matrix
```

```
##      (Intercept) hs_grad home_own I(median_income/1000) party_majRep
## 1             1  76.78209  73.38657             36.96294             1
## 2             1  86.59310  63.22759             59.35103             1
```

```
## 3      1 81.82000 69.65333      42.41793      1
## 4      1 78.86400 71.23067      34.98271      1
## 5      1 82.45862 63.26724      55.26572      0
## 6      1 88.13906 72.14531      50.13584      0
```

```
y <- poverty$poverty; head(y) # n by 1 vector
```

```
## [1] 19.90746 12.29310 19.03333 20.00533 14.22069 12.86094
```

```
XtX <- t(X) %*% X # (p + 1) by (p + 1) matrix
colnames(XtX) <- rownames(XtX) <- NULL
XtX
```

```
##      [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 51.000 4323.498 3657.123 2430.910 30.000
## [2,] 4323.498 367497.623 309942.167 207273.903 2498.765
## [3,] 3657.123 309942.167 264098.108 173256.517 2202.333
## [4,] 2430.910 207273.903 173256.517 119871.817 1284.466
## [5,] 30.000 2498.765 2202.333 1284.466 30.000
```

```
Xty <- t(X) %*% y; head(Xty) # (p + 1) by 1 vector
```

```
##      [,1]
## (Intercept) 734.9980
## hs_grad 61590.5786
## home_own 52725.4769
## I(median_income/1000) 33676.2834
## party_majRep 476.7111
```

```
as.numeric(beta.hat <- solve(XtX, Xty)) # regression coefficients
```

```
## [1] 76.0733587 -0.4534462 -0.1461374 -0.2610123 -0.5109967
```

```
y.hat <- as.numeric(X %*% beta.hat); head(y.hat) # fitted values
```

```
## [1] 20.37351 11.56579 17.21084 20.26140 15.01207 12.47784
```

```
eps.hat <- y - y.hat; head(eps.hat) # residuals
```

```
## [1] -0.4660507 0.7273168 1.8224961 -0.2560688 -0.7913801 0.3830990
```

```
(sigma.hat <- sqrt(sum(eps.hat ^ 2) / (51 - 4 - 1))) # residual SE
```

```
## [1] 1.353753
```

```
fitted(mod) # obtain fitted values
```

```
##      1      2      3      4      5      6      7      8
## 20.373513 11.565787 17.210837 20.261402 15.012070 12.477839 7.194859 11.576235
##      9     10     11     12     13     14     15     16
## 15.217284 15.819316 19.729882 13.879387 14.791221 13.493640 13.587406 11.915343
##     17     18     19     20     21     22     23     24
## 13.474400 20.683055 19.524389 13.240073 8.420077 9.104415 12.953813 11.231420
##     25     26     27     28     29     30     31     32
## 21.637436 17.040754 14.251766 12.859763 13.400985 9.679679 7.710733 17.805235
##     33     34     35     36     37     38     39     40
## 12.902249 17.969372 14.065428 13.752141 16.863705 14.909914 13.171882 9.942183
##     41     42     43     44     45     46     47     48
## 19.186426 15.105553 19.505661 18.487914 10.424302 11.556045 15.513875 14.018749
##     49     50     51
```


```
## 18.759122 11.743215 9.996266
```

```
resid(mod) # obtain residuals
```

```
##          1          2          3          4          5          6
## -0.46605065  0.72731679  1.82249609 -0.25606877 -0.79138008  0.38309895
##          7          8          9         10         11         12
##  1.11764146 -0.07623466  3.28271639  0.17919156  0.05250815 -4.87938690
##         13         14         15         16         17         18
## -0.26394784 -0.55344392 -0.99936285 -0.85271642 -1.16773370 -0.14222118
##         19         20         21         22         23         24
##  0.53654811  0.39742676  0.88408973  1.40272776  2.14377701 -0.33141953
##         25         26         27         28         29         30
##  2.75158873 -0.82944933  0.89109102 -0.96191335 -1.05392611 -0.53967877
##         31         32         33         34         35         36
##  1.15117164  1.54627966  0.03323485 -0.38737151 -1.72769199  0.18081322
##         37         38         39         40         41         42
##  0.17136036  0.16508568 -0.87188216 -1.04218336  0.79835675  1.08687105
##         43         44         45         46         47         48
## -1.23408238 -1.25563045  1.01362932 -0.14890176 -2.07432310  1.12740532
##         49         50         51
## -0.32639451 -0.05154804 -0.56148303
```

```
diag.plots <- lindia::gg_diagnose(mod, theme = theme_bw(), plot.all = FALSE,
                                   scale.factor = 1)
```


```
diag.plots$res_fitted
```



```
diag.plots$residual_hist
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Histogram of Residuals



```
diag.plots$qqplot
```

Normal-QQ Plot



```
(var.beta.hat <- sigma.hat ^ 2 * solve(XtX)) # Estimate of Var(beta.hat)
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 20.995620883 -0.1849707465 -0.0674556619 -0.0029825145 -0.509316515
## [2,] -0.184970747  0.0031353530 -0.0004649709 -0.0010121404  0.001289793
## [3,] -0.067455662 -0.0004649709  0.0012520320  0.0003869122 -0.002294815
## [4,] -0.002982514 -0.0010121404  0.0003869122  0.0011746260  0.008589991
## [5,] -0.509316515  0.0012897933 -0.0022948145  0.0085899906  0.263655017
```

```
(se.beta.hat <- sqrt(diag(var.beta.hat))) # SE beta.hat
```

```
## [1] 4.58209787 0.05599422 0.03538406 0.03427282 0.51347348
```

```
as.numeric(beta.hat / se.beta.hat) # test statistic value
```

```
## [1] 16.6022990 -8.0980881 -4.1300351 -7.6157228 -0.9951764
```

```
as.numeric( # p-values
  pt(abs(beta.hat / se.beta.hat), df = 50, lower.tail = FALSE)
```

```

)

## [1] 2.753456e-22 5.867138e-11 6.883705e-05 3.281345e-10 1.622212e-01
(total.SS <- sum((y - mean(y)) ^ 2)) # Total SS

## [1] 738.7815
(resid.SS <- sum(eps.hat ^ 2)) # Resid SS

## [1] 84.30172
(reg.SS <- total.SS - resid.SS) # Reg SS

## [1] 654.4798
(reg.SS / 4) / (resid.SS / (51 - 4 - 1)) # F-statistic

## [1] 89.28071
1 - resid.SS / total.SS # R^2 value

## [1] 0.8858909
(tmp <- apply(X, 2, mean))

##           (Intercept)           hs_grad           home_own
##           1.0000000           84.7744689           71.7082973
## I(median_income/1000)       party_majRep
##           47.6648981           0.5882353

newx <- data.frame(
  hs_grad = tmp[2],
  home_own = tmp[3],
  median_income = tmp[4] * 1000,
  party_maj = "Dem"
)
predict(mod, newx, interval = "confidence", level = 0.95) # narrow

##           fit           lwr           upr
## hs_grad 14.71231 13.99451 15.43011
predict(mod, newx, interval = "prediction", level = 0.95) # wider

##           fit           lwr           upr
## hs_grad 14.71231 11.89439 17.53023

```

